

# Naïve Bayes Spam Filter

Jenian Tai (kt2694)  
10/7/2018

```
readDirectory <- function(dirname) {  
  # Store the emails in a list  
  emails = list();  
  # Get a list of filenames in the directory  
  filenames = dir(dirname, full.names=TRUE);  
  
  for (i in 1:length(filenames)){  
    emails[[i]] = scan(filenames[i], what="", quiet=TRUE);  
  }  
  return(emails)  
}
```

Loaded the data.

```
ham_test <- readDirectory("/Users/jenian/Documents/APANPS5335/hw4/ham-v-spam/ham-test")  
ham_train <- readDirectory("/Users/jenian/Documents/APANPS5335/hw4/ham-v-spam/ham-train")  
spam_test <- readDirectory("/Users/jenian/Documents/APANPS5335/hw4/ham-v-spam/spam-test")  
spam_train <- readDirectory("/Users/jenian/Documents/APANPS5335/hw4/ham-v-spam/spam-train")
```

```
print(ham_train[1])
```

```
## [[1]]  
## [1] "gari" "product" "high" "island" "larger"  
## [6] "block" "commenc" "saturday" "gross" "carlo"  
## [11] "expect" "gross" "tomorrow" "vastar" "own"  
## [16] "gross" "product" "georg" "daren" "farmer"  
## [21] "carlo" "rodriguez" "hou" "ect" "ect"  
## [26] "georg" "weissman" "hou" "ect" "ect"  
## [31] "melissa" "grave" "hou" "ect" "ect"  
## [36] "carlo" "pleas" "call" "linda" "get"  
## [41] "everyth" "set" "estim" "come" "tomorrow"  
## [46] "increas" "follow" "day" "base" "convers"  
## [51] "bill" "fischer" "bmar" "enron" "north"  
## [56] "america" "corp" "georg" "weissman" "daren"  
## [61] "farmer" "hou" "ect" "ect" "gari"  
## [66] "bryan" "hou" "ect" "ect" "melissa"  
## [71] "grave" "hou" "ect" "ect" "darren"  
## [76] "attach" "appear" "nomin" "vastar" "resourc"  
## [81] "inc" "high" "island" "larger" "block"  
## [86] "previous" "erron" "refer" "well" "vastar"  
## [91] "expect" "well" "commenc" "product" "sometim"  
## [96] "tomorrow" "told" "linda" "harri" "get"  
## [101] "telephon" "number" "gas" "control" "provid"  
## [106] "notif" "turn" "tomorrow" "linda" "number"  
## [111] "record" "voic" "fax" "would" "pleas"
```

```
## [116] "see"      "someone" "contact" "linda"   "advis"
## [121] "submit"   "futur"   "nomin"   "via"     "mail"
## [126] "fax"      "voic"    "thank"   "georg"   "linda"
## [131] "harri"    "georg"   "weissman" "hou"     "ect"
## [136] "ect"      "effect"  "mscf"    "min"     "ftp"
## [141] "time"     "hour"    "hour"    "hour"    "hour"
## [146] "hour"     "hour"    "hour"    "hour"    "hour"
## [151] "hour"     "hour"    "hour"    "hour"    "hour"
## [156] "hour"     "hour"    "hour"    "hour"    "hour"
```

```
print(spam_train[1])
```

```
## [[1]]
##      [1]"introduc"  "doctor"    "formul"    "hgh"       "human"
##      [6]"growth"    "hormon"    "also"      "call"      "hgh"
## [11]  "refer"      "medic"     "scienc"    "master"    "hormon"
## [16]  "plenti"     "young"     "near"      "age"       "twenti"
## [21]  "one"        "bodi"      "begin"     "produc"    "less"
## [26]  "time"       "forti"     "near"      "everyon"   "defici"
## [31]  "hgh"        "eighti"    "product"   "normal"    "diminish"
## [36]  "least"      "advantag"  "hgh"       "increas"   "muscl"
## [41]  "strength"   "loss"      "bodi"      "fat"       "increas"
## [46]  "bone"       "densiti"   "lower"     "blood"     "pressur"
## [51]  "quicken"    "wound"     "heal"      "reduc"     "cellulit"
## [56]  "improv"     "vision"    "wrinkl"    "disappear" "increas"
## [61]  "skin"       "thick"     "textur"    "increas"   "energi"
## [66]  "level"      "improv"    "sleep"     "emot"      "stabil"
## [71]  "improv"     "memori"    "mental"    "alert"     "increas"
## [76]  "sexual"     "potenc"    "resist"    "common"    "ill"
## [81]  "strengthen" "heart"     "muscl"     "control"   "cholesterol"
## [86]  "control"    "mood"      "swing"     "new"       "hair"
## [91]  "growth"     "color"     "restor"    "read"      "websit"
## [96]  "unsubscrib"
```

```
makeSortedDictionaryDf <- function(emails){
  # This returns a dataframe that is sorted by the number of times
  # a word appears
  # List of vectors to one big vector
  dictionaryFull <- unlist(emails)
  # Tabulates the full dictionary
  tabulateDic <- tabulate(factor(dictionaryFull))
  # Find unique values
  dictionary <- unique(dictionaryFull)
  # Sort them alphabetically
  dictionary <- sort(dictionary)
  dictionaryDf <- data.frame(word = dictionary, count = tabulateDic)
  sortDictionaryDf <- dictionaryDf[order(dictionaryDf$count,decreasing=TRUE),]; return(sortDictionaryDf)
}
```

```
all_emails <- c(ham_test, ham_train, spam_test, spam_train)
dictionary <- makeSortedDictionaryDf(all_emails)
```

```
makeDocumentTermMatrix <- function(emails, dictionary){
  # This takes the email and dictionary objects from above and outputs a
  # document term matrix
  num_emails <- length(emails);
  num_words <- length(dictionary$word);
  # Instantiate a matrix where rows are documents and columns are words dtm <-
  mat.or.vec(num_emails, num_words); # A matrix filled with zeros
  for (i in 1:num_emails){
    num_words_email <- length(emails[[i]]);
    email_temp <- emails[[i]];
    for (j in 1:num_words_email){
      ind <- which(dictionary$word == email_temp[j]);
      dtm[i, ind] <- dtm[i, ind] + 1;
    }
  }
  return(dtm);
}
```

```
dtm_ham_train <- makeDocumentTermMatrix(ham_train, dictionary)
dtm_spam_train <- makeDocumentTermMatrix(spam_train, dictionary)
dtm_ham_test <- makeDocumentTermMatrix(ham_test, dictionary)
dtm_spam_test <- makeDocumentTermMatrix(spam_test, dictionary)
```

```
makeLogPvec <- function(dtm, mu){
  # Sum up the number of instances per word
  pvecNoMu <- colSums(dtm)
  # Sum up number of words
  nWords <- sum(pvecNoMu)
  # Get dictionary size
  dicLen <- length(pvecNoMu)
  # Incorporate mu and normalize
  logPvec <- log(pvecNoMu + mu) - log(mu*dicLen + nWords)
  return(logPvec)
}
```

```
mu <- 1 / length(dictionary$word)
log_pvec_ham <- makeLogPvec(dtm_ham_train, mu)
log_pvec_spam <- makeLogPvec(dtm_spam_train, mu)
```

```

predictNaiveBayes <- function(log_pvec_ham, log_pvec_spam, log_ham_prior, log_spam_prior, dtm_test) {
  n <- nrow(dtm_test)
  yh <- numeric(n)
  for (i in 1:n) {
    if (sum(dtm_test[i,] * log_pvec_spam) + log_spam_prior > sum(dtm_test[i,] * log_pvec_ham) + log_ham_prior) {
      yh[i] <- 1
    } else {
      yh[i] <- 0
    }
  }
  return(yh)
}

```

```

ham_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_ham_test)
spam_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_spam_test)

hr <- mean(spam_hat)
fa <- mean(ham_hat)
acc <- (sum(spam_hat) + length(ham_hat) - sum(ham_hat)) / (length(ham_hat) + length(spam_hat))
print(hr)

```

```
## [1] 0.9466667
```

```
print(fa)
```

```
## [1] 0.04
```

```
print(acc)
```

```
## [1] 0.9533333
```

```

fiveFoldCV <- function(dtm_ham_train, dtm_spam_train, log_ham_prior, log_spam_prior, mu){
  errors <- numeric(5)
  # split up your data into 5 sets
  n <- nrow(dtm_ham_train) fold_size <- n/5
  for (i in 1:5) { full_range
    <- 1:n
    validation_range <- ((i-1) * fold_size + 1):(i * fold_size) train_range <-
    full_range[!full_range %in% validation_range]
    # train on the train_range using makeLogPvec()
    log_pvec_ham <- makeLogPvec(dtm_ham_train[train_range,], mu) # 1 point
    log_pvec_spam <- makeLogPvec(dtm_spam_train[train_range,], mu) # 1 point

    # validate on the validation_range using predictNaiveBayes()
    ham_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_ham_train[validation_range,])
    spam_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_spam_train[validation_range,])
    # calculate the error rate and store in vector (did you initialize it?)
    errors[i] <- (mean(ham_hat) + mean(1 - spam_hat))/2 # 1 point
  }
  # return the average error over all folds
  return(mean(errors))
}

```

```

bm <- 1 / length(dictionary$word)
mus <- c(1/100, 1/10, 1, 10, 100) * bm
errs <- numeric(length(mus))
for (i in 1:length(mus)) {
  errs[i] <- fiveFoldCV(dtm_ham_train, dtm_spam_train, log(0.5), log(0.5), mus[i])
}
best_mu <- mus[which.min(errs)]
print(best_mu)

```

```
## [1] 0.0004389623
```

```

log_pvec_ham <- makeLogPvec(dtm_ham_train, best_mu)
log_pvec_spam <- makeLogPvec(dtm_spam_train, best_mu)

ham_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_ham_test)
spam_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_spam_test)

hr <- mean(spam_hat)
fa <- mean(ham_hat)
acc <- (sum(spam_hat) + length(ham_hat) - sum(ham_hat)) / (length(ham_hat) + length(spam_hat))
print(hr)

```

```
## [1] 0.94
```

```
print(fa)
```

```
## [1] 0.04
```

```
print(acc)
```

```
## [1] 0.95
```

The accuracy, specificity and sensitivity turn out lower using the new mu, but we shall stay with new mu =  $1/|D|$  to avoid overfit.

```

calculateMI <- function(dtm_ham_train, dtm_spam_train) {
  # calculates vector of mutual information for each word.
  ham_sums <- colSums(dtm_ham_train)
  ham_probs <- ham_sums / sum(ham_sums) # vector of probabilities for each word in ham
  spam_sums <- colSums(dtm_spam_train)
  spam_probs <- spam_sums / sum(spam_sums) # vector of probabilities for each word in spam
  all_sums <- ham_sums + spam_sums
  all_probs <- all_sums / sum(all_sums) # vector of probabilities for word in entire set mi <-
  c(length(all_probs))
  for (i in 1:length(all_probs)) {
    if (all_probs[i] == 0 || ham_probs[i] == 0 || spam_probs[i] == 0) { mi[i] <- 0
      # mutual information -> 0 when p(X=x) = 0
    }
    else {
      mi[i] <- .5 * ham_probs[i] * log(ham_probs[i] / all_probs[i]) +
        .5 * (1 - ham_probs[i]) * log((1 - ham_probs[i]) / (1 - all_probs[i])) +
        .5 * spam_probs[i] * log(spam_probs[i] / all_probs[i]) +
        .5 * (1 - spam_probs[i]) * log((1 - spam_probs[i]) / (1 - all_probs[i]))
    }
  }
  return(mi)
}

```

```

mivec <- calculateMI(dtm_ham_train, dtm_spam_train)

```

```

fitAndPredict <- function(dtm_ham_train, dtm_spam_train, dtm_ham_test, dtm_spam_test, mu)
{ log_pvec_ham <- makeLogPvec(dtm_ham_train, mu)
  log_pvec_spam <- makeLogPvec(dtm_spam_train, mu)
  ham_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_ham_test)
  spam_hat <- predictNaiveBayes(log_pvec_ham, log_pvec_spam, log(.5), log(.5), dtm_spam_test)
  hr <- mean(spam_hat)
  fa <- mean(ham_hat)
  acc <- (sum(spam_hat) + length(ham_hat) - sum(ham_hat)) / (length(ham_hat) + length(spam_hat))
  return(c(hr, fa, acc))
}

```

```

ns <- c(200,500,1000,2500,5000,10000)
hrs <- c(length(ns))
fas <- c(length(ns))
accs <- c(length(ns))
for(i in 1:length(ns)) {
  n <- ns[i]

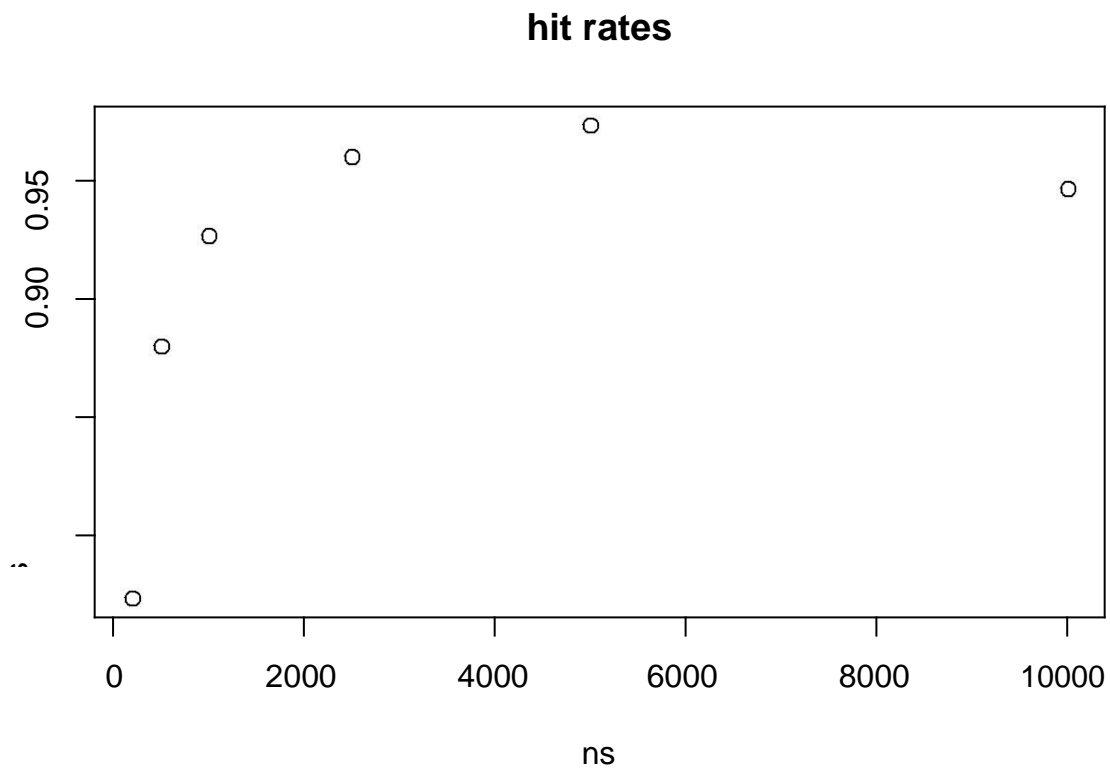
```

```

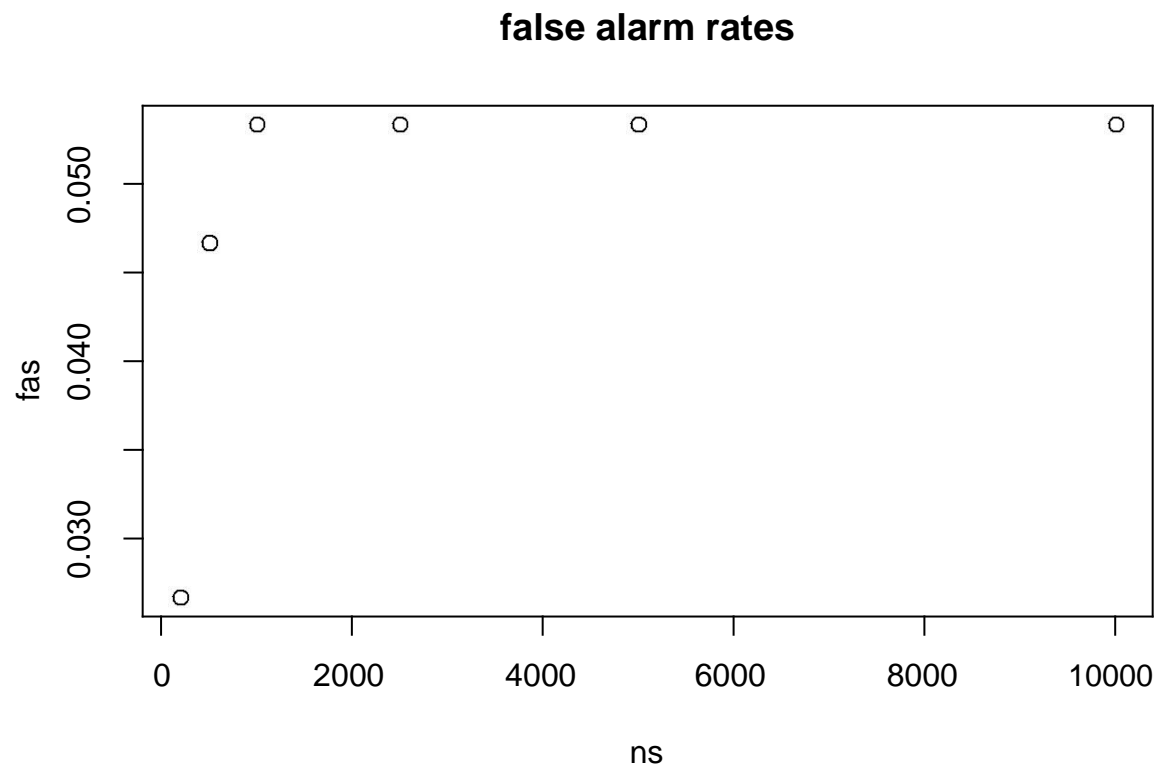
sub_dtm_ham_train <- dtm_ham_train[,order(mivec)[1:n]]
sub_dtm_spam_train <- dtm_spam_train[,order(mivec)[1:n]]
sub_dtm_ham_test <- dtm_ham_test[,order(mivec)[1:n]]
sub_dtm_spam_test <- dtm_spam_test[,order(mivec)[1:n]]
o <- fitAndPredict(sub_dtm_ham_train, sub_dtm_spam_train,
                   sub_dtm_ham_test, sub_dtm_spam_test, 1/n)

hrs[i] <- o[1]
fas[i] <- o[2]
accs[i] <- o[3]
}
plot(ns, hrs, main="hit rates")

```



```
plot(ns, fas, main="false alarm rates")
```



```
plot(ns, accs, main="accuracy")
```

